Market Regime Detection Using Gaussian Hidden Markov Models

Tumelo Tshwanelo Mkwambe

School of Computer Science and Applied Mathematics
University of the Witwatersrand
Johannesburg, South Africa
2446873@students.wits.ac.za

Abstract—A market regime refers to a period during which financial markets exhibit consistent and persistent behaviors or patterns. The most common regimes are directional (trend-based), such as bull markets, characterized by investor optimism and rising index levels, and bear markets, marked by investor pessimism and declining prices. Another key feature that defines market regimes is volatility—a statistical measure of the degree of variation in the price of a financial instrument over time, often used as a proxy for market uncertainty or risk. In this study, we apply Gaussian Hidden Markov Models (HMMs) to identify and characterize distinct market regimes in the South African Top 40 Index.

I. INTRODUCTION

The South African Top 40 (JTOPI) index is a capitalization-weighted index that represents the performance of the top 40 companies listed in the Johannesburg Stock Exchange (JSE). The companies are listed by their free-float market capitalization, which reflects shares that are publicly available for trading. The composition of the South African Top 40 is reviewed periodically to ensure it accurately reflects the largest companies on the JSE. The index is considered a market proxy for the South African stock market.

Hidden Markov Models (HMMs) are a class of dynamic Bayesian networks designed to model sequential data. An HMM is composed of two sets of variables or states, hidden/latent and observed. The hidden states represent the underlying conditions of a system that evolves over time. The observed variables are the data we directly observe; their distribution depends only on the current hidden state.

Hidden Markov Models (HMMs) view the market as switching between hidden states, each exhibiting unique patterns of returns, volatility, and other characteristics. By analyzing market data, HMMs try to uncover these hidden states and understand how the market transitions between them. This makes HMMs very competent in detecting market regimes compared to other statistical methods.

II. PROBLEM FORMULATION

Prior to constructing and evaluating investment strategies and assessing risk, it is required for quantitative analysts to be aware of the current market conditions. A substantial part of this consists of knowing investors' attitudes towards a certain market, and the current underlying sentiment can be analyzed from the trading price. Market sentiment is shaped by numerous factors, including political events, company performance, and environmental occurrences, as exemplified by the Covid-19 outbreak triggering a significant bear trend with high volatility across many markets. However, market sentiment itself represents a collective quantification of investors' expectations regarding the impact of these factors. Therefore, the problem is to develop a robust and timely methodology to infer prevailing market sentiment from historical and real-time trading data.

III. DATA AND ANALYSIS

A. Data

The historical time series data for the South African top 40 index was downloaded from Investing.com in a csv format. Investing.com is a financial markets platform and news website. The downloaded csv file contains daily close price, open price and volume data of the index from 2015-01-01 to 2025-01-01. The entries in the csv file are in string format and contained 13 null entries for the Volume column. Post float-casting and data cleaning, the data was well prepared for analysis.

B. Analysis

Raw prices, while representing the direct value of an asset over time, suffer from several limitations that make them less suitable for many financial analyses. The traditional approach is to use log returns instead, which are mathematically convenient. Log returns are additive over time and exhibit stationarity (a property that raw prices lack) and approximate the Gaussian probability distribution (which complements the Gaussian emission assumption for Gaussian hidden Markov models).

The Stationarity / Time Invariance property for hidden Markov models states that the probability of transitioning between states remains constant.

$$P(X(t+1) = \xi | X(t) = \xi) = P(X = \xi | X = \xi)$$

The stationarity assumption in HMMs assumes the process is stationary. This justifies the use of log returns in constructing the model. Do note that time series data is stationary if its statistical properties (mean, variance, etc.) remain constant.

Fig. 1. Rolling mean and standard deviation for closing price.

Fig. 2. Rolling mean and standard deviation for log returns.

Fig. 1 illustrates the non-stationarity of the index's closing price. In contrast, Fig. 2 demonstrates that the log returns exhibit stationarity, with a notable spike in the middle due to the COVID-19 outbreak. Furthermore, the suitability of log returns for training Gaussian hidden Markov models is supported by their approximate normal distribution, as shown in Fig. 3. For Gaussian hidden Markov models, the emission probability associated with each hidden state is modeled as a normal distribution.


Fig. 3. Log Return Distribution.

The summary statistics for the log returns of the index are displayed in table I.

IV. HIDDEN MARKOV MODELS

A Markov process is a stochastic process whereby the next state only depends on the current state and is independent of all other previous states.

Let $X_0, X_1, ..., X_{t-1}, X_t$ be a sequence of random variables or states of a stochastic process, then the Markov assumption states that:

$$P(X_t = x | X_0, ..., X_{t-1}) = P(X_t = x | X_{t-1}). \quad (1)$$

A hidden Markov model is a Markov process specified by the following components:

- **States:** $Q = \{q_0, q_1, ..., q_{N-1}\}$
- **Transition Probability Matrix:** $A \in \mathbb{R}^{N \times N}$, whereby $a_{i,j}$ is the probability of transitioning from state i into state j.
- **Emission Probability Matrix**: $B \in \mathbb{R}^{N \times M}$ for an observation alphabet $O \in \mathbb{R}^M$, whereby $b_{ij} = b_i(O_t = O_j) = b_i(O_j)$ is the probability of observation symbol O_j at time t given state q_j at time t.
- **Initial Probabilities:** $\pi \in \mathbb{R}^N$ where π_i is the probability of being in state i at time t_0,

where N is the number of states.

Hidden Markov models can be used to address three core problems. In this section, we provide a brief overview of these problems, followed by efficient algorithms for solving them.

A. The Likelihood Problem and the Forward Algorithm

The likelihood problem is stated as follows: Given a hidden Markov model $\lambda = (A, B, \pi)$ and a sequence of observations $O = \{O_0, O_1, \ldots, O_{T-1}\}$, what is the probability of the observation sequence given the model, $P(O|\lambda)$?

To compute $P(O|\lambda)$, we use the *forward algorithm*, which recursively calculates the probability of observing the sequence up to time t and being in state q_i at time t. Let

$$\alpha_t(i) = P(O_0, O_1, \ldots, O_t, q_t = q_i \mid \lambda),$$

then the algorithm proceeds as follows:

$$\alpha_0(i) = \pi_i b_i(O_0), \qquad \text{for } i = 0, 1, \ldots, N-1, \tag{2}$$

$$\alpha_t(i) = b_i(O_t) \sum_{j=0}^{N-1} \alpha_{t-1}(j) a_{ji}, \quad \text{for } t = 1, \ldots, T-1, \tag{3}$$

$$P(O|\lambda) = \sum_{i=0}^{N-1} \alpha_{T-1}(i). \tag{4}$$

B. Posterior State Estimation and the Backward Algorithm

Given a model $\lambda = (A, B, \pi)$ and an observation sequence $O = \{O_0, O_1, \ldots, O_{T-1}\}$, we wish to compute the posterior probability of being in state q_i at time t, given the full observation sequence:

$$P(q_t = q_i \mid O, \lambda).$$

This problem is commonly referred to as *soft-decoding*, in contrast to *hard-decoding*, which aims to find the most likely sequence of states. While hard-decoding is solved via the Viterbi algorithm, soft-decoding makes use of the *backward algorithm*. Define:

$$\beta_t(i) = P(O_{t+1}, O_{t+2}, \ldots, O_{T-1} \mid q_t = q_i, \lambda),$$

then the backward algorithm computes:

$$\beta_{T-1}(i) = 1, \tag{5}$$

$$\beta_t(i) = \sum_{j=0}^{N-1} a_{ij} b_j(O_{t+1}) \beta_{t+1}(j), \quad \text{for } t = T-2, \ldots, 0. \tag{6}$$

The posterior probability of being in state q_i at time t is:

$$\gamma_t(i) = P(q_t = q_i \mid O, \lambda) = \frac{\alpha_t(i)\beta_t(i)}{P(O|\lambda)}. \tag{7}$$

The most probable state at time t is the one maximizing $\gamma_t(i)$ over all i.

C. Parameter Estimation and the Baum-Welch Algorithm

The Baum-Welch algorithm is a variant of the Expectation-Maximization (EM) algorithm that estimates the parameters of a hidden Markov model from data. It proceeds in two main steps:

E-step: Compute the expected values of hidden variables using the current parameter estimates: - Forward probabilities $\alpha_t(i)$ - Backward probabilities $\beta_t(i)$ - Posterior state probabilities $\gamma_t(i)$ - Posterior transition probabilities $\xi_t(i, j)$, defined as:

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{\sum_{k=0}^{N-1} \sum_{l=0}^{N-1} \alpha_t(k) a_{kl} b_l(O_{t+1}) \beta_{t+1}(l)}. \tag{8}$$

M-step: Update the model parameters using the computed expectations:

- **Initial state distribution:**

$$\pi_i = \gamma_0(i)$$

- **Transition probabilities:**

$$A_{ij} = \frac{\sum_{t=0}^{T-2} \xi_t(i, j)}{\sum_{t=0}^{T-2} \gamma_t(i)}$$

- **Emission probabilities (discrete case):**

$$b_i(k) = \frac{\sum_{t=0}^{T-1} \mathbb{I}[O_t = O_k] \gamma_t(i)}{\sum_{t=0}^{T-1} \gamma_t(i)},$$

where $\mathbb{I}[\cdot]$ is the indicator function.

D. Gaussian Emission Probabilities

In a Gaussian hidden Markov model, emissions are continuous and modeled by Gaussian distributions. Each hidden state q_i emits observations $O_t \in \mathbb{R}$ according to:

$$P(O_t = o \mid q_t = q_i) = \mathcal{N}(o; \mu_i, \sigma_i^2).$$

The parameters of these distributions are updated during the M-step as follows:

$$\mu_i = \frac{\sum_{t=0}^{T-1} \gamma_t(i) O_t}{\sum_{t=0}^{T-1} \gamma_t(i)}, \tag{9}$$

$$\sigma_i^2 = \frac{\sum_{t=0}^{T-1} \gamma_t(i)(O_t - \mu_i)^2}{\sum_{t=0}^{T-1} \gamma_t(i)}. \tag{10}$$

E. Gaussian Emission Probabilities

When using a Gaussian hidden Markov model, the emission probabilities are not discrete, but instead modeled using Gaussian distributions. Each hidden state emits an observation $O_j \in \mathbb{R}$ in accordance with its distribution: $P(O_t = O_j | q_t = q_i) = \mathcal{N}(O_j; \mu, \sigma^2)$. Consequently, the parameters of the distribution in each state are updated during the M-step as:

$$\mu_i = \frac{\sum_{t=0}^{T-1} \gamma_t(i) O_t}{\sum_{t=0}^{T-1} \gamma_t(i)}, \tag{11}$$

$$\sigma_i^2 = \frac{\sum_{t=0}^{T-1} \gamma_t(i)(O_t - \mu_i)^2}{\sum_{t=0}^{T-1} \gamma_t(i)}. \tag{12}$$

F. Hard-decoding And The Viterbi Algorithm

Hard-decoding refers to the process of finding the single most probable sequence of hidden states given a sequence of observations. Let $\delta_t(i)$ denote the highest probability of any state sequence that ends in state q_i at time t, given the observation sequence up to time t. The Viterbi algorithm computes $\delta_t(i)$ recursively for each state i at each time step t, using the previous values $\delta_t(j)$. At each step, it selects the transition from state q_j that maximizes the probability of reaching q_i, and it stores a pointer $\psi_t(i)$ to the state q_j that

yields this maximum. This allows the most probable state sequence to be recovered through backtracking after the final time step:

Initialization:

$$\delta_0(i) = \pi_i b_i(O_0) \tag{13}$$

$$\psi_0(i) = 0 \tag{14}$$

Iteration:

$$\delta_t(i) = b_i(O_t) \max_k [\delta_{t-1}(k) a_{ki}] \tag{15}$$

$$\psi_t(i) = \max_k [\delta_{t-1}(k) a_{ki}] \tag{16}$$

Termination:

$$P^* = \max_i \delta_{T-1}(i) \tag{17}$$

$$Q^* = \arg\max_i \delta_t(i) \tag{18}$$

Backtracking:

$$Q_t^* = \psi_{t+1}(Q_{t+1}^*), \text{ for } t = T-2, T-3, ..., 0 \tag{19}$$

V. MODEL IMPLEMENTATION AND RESULTS

We use the Baum-Welch algorithm to train two different models on the same South African top 40 index dataset. The first model has two distinct states and the other three. The initial parameters for both models are randomized prior to training. Each model is allowed to execute the Baum-Welch algorithm for a maximum of 100 iterations given that the change in log-likelihood is no greater than a threshold of 0.00001 for convergence. We make use of log probabilities and appropriate scaling to avoid underflow during the Viterbi, forward and backward algorithms.

A. Two-State Model Specifications And Results

The estimated parameters post-training for the two-state model are presented in Fig. 4.

	STATE 0	STATE 1
INITIAL PROBABILITY	1.0	0.0
MEAN	-0.0002814	0.0003763
VARIANCE	0.0003934	0.00007779
FROM STATE 0	0.92904842	0.07095158
FROME STATE 1	0.01779979	0.98220021

Fig. 4. Parameters for two-state model.

State 0 represents a regime of high volatility in a bear market, this is signaled by the high positive variance and negative mean. State 1 represents steady sideway and bullish regimes of low volatility, signaled by the positive mean (close to the log return mean) and low variance.
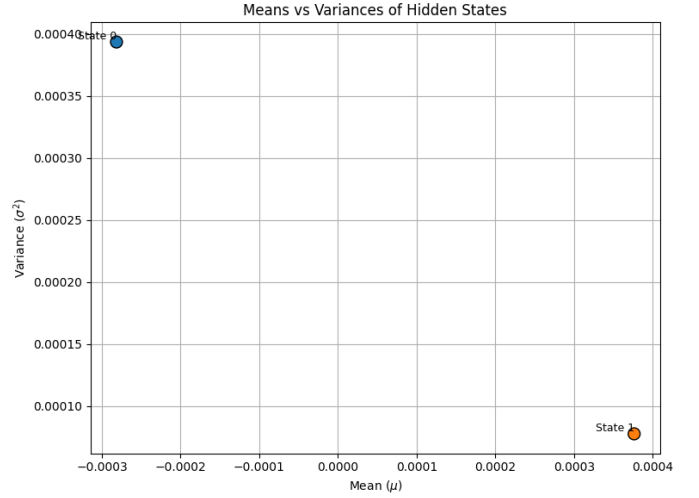

Fig. 5. Mean-variance plot for the two-state model.

The model is capable of classifying regimes of high and low volatility, exemplified by the detection of the covid-19 outbreak as a bearish period of very high volatility as seen in Fig 6.


Fig. 6. Two-state model market regime detection.

B. Three-State Model Specifications And Results

The estimated parameters for the three-state model are presented in Fig. 7.

	STATE 0	STATE 1	STATE 2
INITIAL PROBABILITY	0.0	1.0	0.0
MEAN	0.00066296	-0.0044601337	-0.000229067
VARIANCE	0.00006.12265	0.00166678	0.0002054099
FROM STATE 0	0.963204310	0.000745900914	0.0360497892
FROME STATE 1	0.0	0.947350845	0.0526491553
FROM STATE 3	0.0596516715	0.0	0.940348328

Fig. 7. Parameters for three-state model.

State 0 displays a low positive mean and low variance, which is an indication of a steady bull trend of low volatility.

State 1 represents a period of crisis, this is due to the high volatile bear trend, as seen in early 2020, in contrast state 2 represents a steady bear trend of low volatility.
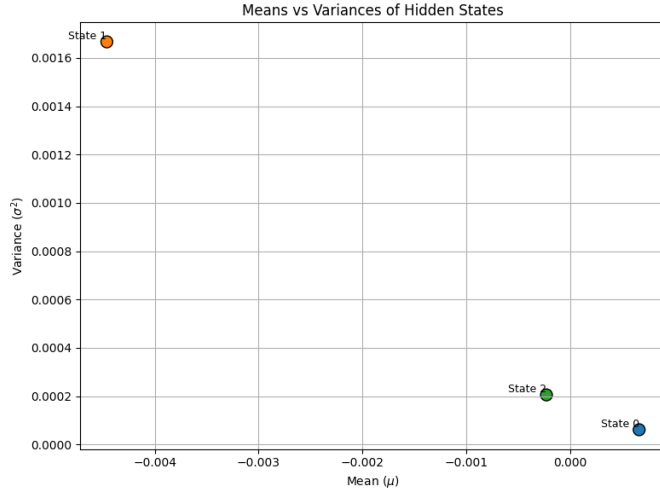

Fig. 8. Mean-variance plot for the three-state model.

Such a model could be leveraged to detect crisis-inspired regimes.

Fig. 9. Three-state model market regime detection.

VI. CONCLUSION

This study demonstrates the effectiveness of Gaussian Hidden Markov Models (HMMs) in detecting and interpreting market regimes within the South African Top 40 Index. By modeling financial time series data as a sequence of hidden states—each representing a distinct market condition such as bull, bear, or crisis periods—HMMs provide a probabilistic framework for understanding how markets evolve over time.

The use of log returns ensured the stationarity of the input data, aligning well with the Gaussian emission assumption of the models. The results from both the two-state and three-state models successfully identified key market regimes, including the high-volatility crisis induced by the COVID-19 pandemic. The two-state model distinguished between high- and low-volatility regimes, while the three-state model further segmented market behavior into nuanced trends, including crisis-specific states.

Overall, the study confirms that Gaussian HMMs are a valuable tool for market regime detection, offering both interpretability and analytical rigor. These models can support better-informed investment strategies, improved risk management, and more timely recognition of market transitions.

VII. REPOSITORY AND CODE

The complete codebase is available in the following GitHub repository: `Gaussian-Hidden-Markov-Model`.

The repository includes the following files:
- `2015_2025_Historical_Data.csv` – Historical market data retrieved from Investing.com.
- `data.py` – Defines a class containing methods for data preprocessing.
- `analysis.py` – Provides a class for time series analysis applicable beyond the JTOPI index.
- `hidden_markov_model.py` – Implements the Hidden Markov Model, including the Viterbi, forward, backward, and Baum-Welch algorithms.
- `main.ipynb` – Jupyter notebook demonstrating data loading, preprocessing, analysis, and HMM training on JTOPI data.
- `two_state_model.pkl` – Serialized parameters for the trained two-state HMM.
- `three_state_model.pkl` – Serialized parameters for the trained three-state HMM.